

NANYANG TECHNOLOGICAL UNIVERSITY
SINGAPORE

Evaluation Report

Five measures and 34 test cases, baseline and LLM mode

AI-driven Project Management System for the Harbourline Community Hub Refurbishment

Course: MA6094 Agentic AI for Project Management
Capstone Project Group 5

Tian Yujia	G2510505D	team lead
Qu Chengzhi	G2510533D	
Zhao Lingfeng	G2510595H	
Zhang Yinan	G2505943K	
Liu Beiqi	G2510916G	
Huang Hao	G2510626K	

Document owner: Qu Chengzhi

Draft for team review

14 November 2026

Abstract

This report evaluates the Harbourline AI project management system on the latest code. The harness `eval/run_eval.py` runs 34 test cases, T01–T34, each in a throw-away database against a mock project board, and computes five measures: gap recall, invention rate, re-run stability, Critic catch rate, and cost and time per full run. Both modes were run on 14 November 2026. In the deterministic baseline mode all 34 cases pass: gap recall is 91%, 20 of 22 gold brief gaps; the invention rate is 0.0%, 0 of 75 sampled claims; re-runs are identical with no duplicates; the Critic stops all 21 planted faults; and a full run takes 0.055 s at no cost. In LLM mode, with DeepSeek, all 34 cases also pass: gap recall is 82%, 18 of 22; the invention rate is 0.0%, 0 of 92; the catch rate is 100%; and a full run takes 92.5 s, 79,812 tokens and US$0.026. LLM outputs are not identical across fresh databases, but re-runs on the same database are, because responses are cached. Gap recall in LLM mode varies between runs, from 77% to 82%. The first LLM-mode runs exposed real defects, among them untraceable figures in generated risks, a missing conflict in the change analysis and a memory question wrongly answered NOT FOUND; these were fixed, and six checks that pinned baseline IDs were rewritten to test behaviour. The report closes with the limitations of the evaluation and the commands to reproduce it.

Contents

Abstract **i**

1 What was evaluated and how **1**
1.1 Isolation . . . 1
1.2 Fault injection and ground truth . . . 1
1.3 Two modes . . . 1
1.4 Tests check behaviour, not baseline IDs . . . 2

2 Results of the five measures **3**
2.1 Overall . . . 3
2.2 Gap recall in LLM mode varies between runs . . . 3

3 Test cases **5**

4 Detail **8**
4.1 Brief gaps: gold list and recall . . . 8
4.2 Critic: planted faults . . . 10
4.3 Week 12 change: what the system produced for CE-01 . . . 10

5 Defects found by the evaluation and what changed **12**
5.1 Fresh-database stability . . . 12

6 Limitations of this evaluation **14**

7 How to reproduce **15**

References **16**

A Per-case details of the baseline run **17**
A.1 Five measures as recorded . . . 17
A.2 Per-case details . . . 17

B Per-case details of the LLM-mode run **21**
B.1 Five measures as recorded . . . 21
B.2 Per-case details . . . 21

C Reproduction commands **24**

1 What was evaluated and how

This report covers the MA6094 Capstone, Group 5, as of 14 November 2026. It is a draft for team review. Team: Group 5, no team name. Evaluation owner: Qu Chengzhi, G2510533D.

The evaluation harness is `eval/run_eval.py`. It has 34 test cases, T01–T34, and computes five measures, set out in Chapter 2. All results in this report come from the latest code on branch `main`. Quantifying such properties of an AI system before it is relied on corresponds to the Measure function of the NIST AI Risk Management Framework (Tabassi, 2023).

1.1 Isolation

- Each run uses a throw-away SQLite database, so the demo database is never touched.
- The PM tool is always the local mock board, which implements the same data contract as the Jira adapter. Jira is used only with `--jira`, and that was not used here.

1.2 Fault injection and ground truth

Planted Critic faults come from `eval/faults.py`. Brief-gap ground truth is `eval/gold_gaps.json`: 22 gaps the team identified in the brief.

1.3 Two modes

Both modes were run on 14 November 2026; Table 1-1 summarises them.

Table 1-1 The two evaluation modes

Mode	What generates the plan, register, reports and change analysis	Command	Evidence
Baseline, deterministic	Curated, hand-checked tables; no LLM calls	`python3 eval/run_eval.py --no-llm`	`eval/results/latest.md`, `latest.json`
LLM, system default	DeepSeek `deepseek-chat` at temperature 0, validated and checked by the Critic	`python3 eval/run_eval.py`	`docs/evaluation/evidence/eval_llm_deepseek_2026-11-14_final.log`, `llm_results_2026-11-14.json` and `.md`

Gemini is configured as the first provider but has no key on the evaluation machine, so every LLM call went to DeepSeek. The course instructor has accepted the use of DeepSeek, according to the team.

1.4 Tests check behaviour, not baseline IDs

The first LLM-mode evaluation on 8 November 2026 passed 27 of 33 cases. Most failures came from checks that expected the IDs or counts of the curated baseline, such as WBS element `3.2.3`, risk R22, re-score targets R06 and R17, or exactly four flagged WBS candidates. A generated plan numbers its items differently. On 14 November 2026 these checks were rewritten to test the behaviour instead, and T34 was added for a memory defect found in review. Chapter 5 lists each change.

2 Results of the five measures

Table 2-1 gives the definition, target and result of each measure in both modes; Figure 2-1 compares the percentage measures with their targets.

Table 2-1 The five measures: definition, target and results in both modes

#	Measure	Definition	Target	Baseline	LLM mode
1	**Gap recall**	Share of the 22 gold brief gaps that appear unprompted in the Sponsor Question register or the conflict register, matched by keyword groups	$\geq 80\%$	**91%**, 20 of 22	**82%**, 18 of 22
2	**Invention rate**	Share of sampled stakeholder-facing claims in reports, risk register and change note that the Critic tracer cannot trace to their cited sources	$\leq 5\%$	**0.0%**, 0 of 75	**0.0%**, 0 of 92
3	**Re-run stability**	a) two fresh databases give identical plan, register, gaps and reports; b) a second PM sync creates and updates nothing; c) no duplicates	identical; 0 / 0 / 0	a) identical on all four; b) 0 created, 0 updated, 92 unchanged; c) none	a) **not identical** across fresh databases; b) 0 created, 0 updated, 111 unchanged; c) none. Re-runs on the same database are identical, T33
4	**Critic catch rate**	Share of planted faults that end `critic_stopped`	$\geq 90\%$	**100%**, 21 of 21	**100%**, 21 of 21
5	**Cost and time per full run**	Mean of two full runs on fresh databases: seconds, tokens, US$	Free route available	**0.055 s, 0 tokens, US$0**	**92.5 s, 79,812 tokens, US$0.026**; a re-run on the same database costs 0 tokens because responses are cached

2.1 Overall

Baseline: 34 of 34 pass. LLM mode: 34 of 34 pass.

2.2 Gap recall in LLM mode varies between runs

Three LLM-mode runs on 14 November 2026 in which the plan was generated gave 82%, 77% and 82%. The 77% run is `eval_llm_deepseek_2026-11-14_run3.log`. The 80% target is the team's own; the course instructor later confirmed a lower threshold, which every run meets, according to the team. Design limitation L1 records this.

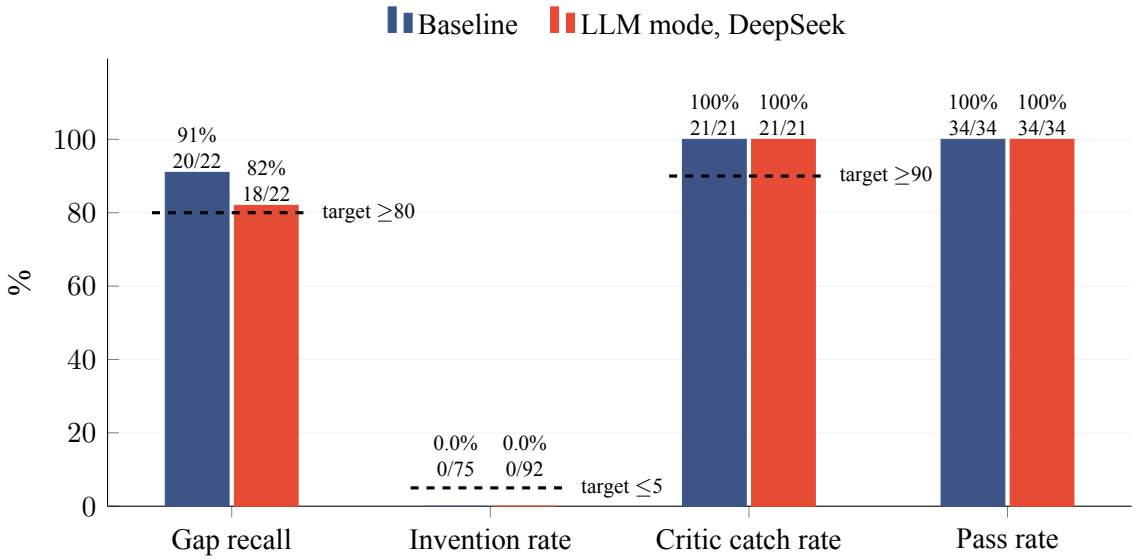


Figure 2-1 Percentage measures in both modes, 14 November 2026 runs. Black dashes mark the targets: at least 80% gap recall, at most 5% invention rate and at least 90% Critic catch rate; the pass rate has no target

3 Test cases

Table 3-1 lists all 34 test cases. "Input" is what the harness does; "Pass if" is the assertion in `eval/run_eval.py`. Cases marked † were rewritten on 14 November 2026 to test behaviour instead of baseline IDs; T34 is new. Figure 3-1 summarises the outcomes of both modes at a glance. The full per-case detail recorded by the harness is reproduced in Appendix A for the baseline run and in Appendix B for the LLM-mode run.



Figure 3-1 Pass and fail matrix of T01–T34 in both modes, 14 November 2026. All 68 results pass. Daggers mark the six cases rewritten as behaviour checks; T34 is the new memory case

Table 3-1 The 34 test cases with their results in baseline mode and LLM mode

ID	Area	Input	Pass if	Baseline	LLM
T01	**Brief gaps**	Full run; match 22 gold gaps against the register text	recall ≥ 0.8	PASS 20/22	PASS 18/22
T02	Planning: conflicts	Full run	CF01, CF02, CF03, CF05, CF06 present, each citing ≥ 2 sources	PASS, CF01–CF12	PASS, CF01–CF12
T03†	Planning: traceability	Full run	No kept WBS element without an in-scope line; at least one untraceable candidate flagged and not kept	PASS, 4 flagged	PASS, 3 flagged
T04†	Planning: WBS	Full run	≥ 40 Level 3 elements with dictionary; every missing owner shown as UNASSIGNED with its reason	PASS, 49 L3	PASS, 61 L3
T05	Planning: milestones	Full run	No dependency issues; undated milestones are marked NOT FOUND or relative, never given an invented date	PASS, 7 undated	PASS, 14 undated
T06	Memory: citation	"What is the approved budget?"	Cites BRIEF-p3-5 or EDC-1.2	PASS	PASS

continued on next page

Table 3-1 continued

ID	Area	Input	Pass if	Baseline	LLM
T07	Memory: citation	"What is the maximum grant?"	Cites NRF-1.2	PASS	PASS
T08	Memory: citation	"How much is the asbestos removal estimate?"	Cites EDC-5.1 or PR3-3	PASS	PASS
T09	Memory: citation	"Who appoints the active ageing centre operator?"	Cites SA-Q2	PASS	PASS
T10	Memory: NOT FOUND	"What is the energy performance target?"	NOT FOUND	PASS	PASS
T11	Memory: NOT FOUND	"How much will the rooftop garden cost?"	NOT FOUND	PASS	PASS
T12	Memory: NOT FOUND	"What is the contractor's insurance policy number?"	NOT FOUND	PASS	PASS
T13	Memory: NOT FOUND	"What is the resident satisfaction target?"	NOT FOUND	PASS	PASS
T14	Memory: conflicts	"When is practical completion?"	Names conflict CF01 and cites ≥ 2 sources	PASS	PASS
T15	Memory: decisions	Record decision DL-001, rebuild the index as a new session, ask about it	Answer cites DL-001	PASS	PASS
T16	Risk: register	Full run	Every risk has L$\times$I = score, owner, trigger, WBS link, source; seed QA finds the R06 score error, R08 closed and the duplicate	PASS, 22 risks	PASS, 24 risks
T17	Risk: what-if W1	Full run	Headroom left 0.58–0.83; completion forecast 2029-01-28 to 2029-02-11; ≥ 4 assumptions	PASS	PASS
T18	Reporting	Full run	EDC and NRF drafts both `critic_passed`	PASS	PASS, LLM-written drafts
T19	Critic	21 planted faults, Section 4.2	$\geq$ 90% stopped	PASS 21/21	PASS 21/21
T20	Governance: HC5	Release a draft with an invented figure	Release refused	PASS	PASS
T21	Tool integration	Preview and apply without HC2 or HC3	No write; refusals returned	PASS	PASS
T22	Re-run: PM sync	Approve and sync twice	Second sync 0 created, 0 updated; no duplicates; read-back verified	PASS, 92 items	PASS, 111 items
T23	Governance: guardrail	Agent tries delete, close, Done and a status update	4 of 4 blocked and logged	PASS	PASS

continued on next page

Table 3-1 continued

ID	Area	Input	Pass if	Baseline	LLM
T24†	**Week 12 change: CE-01**	Inject CE-01, the grant eligibility ruling	W1, at least three risks and a report marked stale; CF12 in UNRESOLVED; ≥ 3 decisions; pending HC6 actions; every cost line traced	PASS	PASS
T25	**Week 12 change: refresh**	After CE-01	New report versions cite CE-01 and pass the Critic, with no manual re-run	PASS	PASS
T26†	**Week 12 change: HC6**	Approve the first CE-01 re-score action	The risk keeps its score while the action is pending, takes the proposed score after approval, and the refreshed reports pass	PASS, R01 20 $\rightarrow$ 25	PASS, R01 20 $\rightarrow$ 25
T27†	**Week 12 change: CE-02**	Inject CE-02, the operator withdrawal	Re-scores proposed for existing risks; affected and downstream milestones found	PASS, R06, R17, R21	PASS, R06, R15, R26
T28	**Week 12 change: Critic**	Post-event draft that ignores CE-01	`critic_stopped` by rule C5	PASS	PASS
T29	Governance: bias	Bias check on 214 consultation responses	65+ = 25 responses; raw and re-weighted top four differ; English-only loses senior comments; status awaiting HC7	PASS	PASS
T30	Governance: audit	After sync	PM write rows $\geq$ issues; no approval without an account	PASS, 92 / 92	PASS, 111 / 111
T31	Integration	Full run and change	Hand-offs planning$\rightarrow$risk, risk$\rightarrow$reporting, reporting$\rightarrow$critic, change$\rightarrow$reporting logged	PASS	PASS
T32	Invention rate	Sample all checkable claims	$\leq 5\%$ untraceable	PASS 0/75	PASS 0/92
T33†	Re-run stability	Re-run the full pipeline twice on the same database, in the same mode	Gap register hash unchanged; no new report versions	PASS	PASS
T34	Memory: frequency	"How often does the EDC meet?"	Answer cites EDC-1.6 or EDC-2.3	PASS	PASS

Brief gaps: T01, with T02 and T03. **Week 12 change:** T24–T28.

4 Detail

4.1 Brief gaps: gold list and recall

Table 4-1 shows, for each of the 22 gold gaps, whether it was surfaced in each mode, with ✓ for surfaced and ✗ for missed; Figure 4-1 shows the same coverage graphically.

Table 4-1 Gold brief gaps G01–G22 and whether each mode surfaced them

Gold gap	Baseline	LLM	Gold gap	Baseline	LLM
G01 Start date / what 22 months covers	✓	✓	G12 Energy performance target	✓	✓
G02 22 months vs "fully open by Q4 2028"	✓	✓	G13 Barrier-free access while lifts replaced	✓	✓
G03 Hub open vs lift works; what may close	✓	✓	G14 Existing leases during works	✓	✗
G04 Who appoints the AAC operator, when	✓	✓	G15 Variation approval authority	✓	✓
G05 Childcare month 18, licensing	✓	✓	G16 Weekly update content and cadence	✓	✓
G06 Grant conditions "will follow"	✓	✓	G17 Rooftop garden feasibility / funding	✓	✓
G07 Who funds the remaining 40%	✓	✓	G18 Hazardous materials in a 1994 building	✓	✓
G08 Is there a contingency	✓	✓	G19 Who operates the Hub after completion	✓	✗
G09 Cost breakdown by element	✓	✓	G20 Are ~100 childcare places fixed	✗	✗
G10 Consultation timing and owner	✓	✓	G21 Working hours / noise limits	✗	✗
G11 Measuring "high resident satisfaction"	✓	✓	G22 60% of S$18.5M vs eligible-cost basis	✓	✓

How matching works. Recall is measured by keyword groups, so a question phrased with different words counts as a miss.

Review of the gold list. On 14 November 2026 each of the 22 gold gaps was checked against the brief's scenario text, lines 77–111 of `data_pack/00_Brief_UPDATED.txt`. Every gap traces to a sentence there: for example G20 to "about 100 children", G21 to "must remain open to residents throughout the works", and G19 to "becomes the estate's flagship shared space" with no operator named. No gap was removed or added. The review was made with Claude and is owned by Qu Chengzhi.

Why both modes miss G20 and G21. Neither the curated checklist nor the generated register raises them. Which other gaps the LLM misses changes from run to run: G14 and G19 in the final run, G12, G15 and G19 in run 3.

	Baseline	LLM mode
G01 Start date / what 22 months covers	●	●
G02 22 months vs "fully open by Q4 2028"	●	●
G03 Hub open vs lift works; what may close	●	●
G04 Who appoints the AAC operator, when	●	●
G05 Childcare month 18, licensing	●	●
G06 Grant conditions "will follow"	●	●
G07 Who funds the remaining 40%	●	●
G08 Is there a contingency	●	●
G09 Cost breakdown by element	●	●
G10 Consultation timing and owner	●	●
G11 Measuring "high resident satisfaction"	●	●
G12 Energy performance target	●	●
G13 Barrier-free access while lifts replaced	●	●
G14 Existing leases during works	●	✕
G15 Variation approval authority	●	●
G16 Weekly update content and cadence	●	●
G17 Rooftop garden feasibility / funding	●	●
G18 Hazardous materials in a 1994 building	●	●
G19 Who operates the Hub after completion	●	✕
G20 Are ~100 childcare places fixed	✕	✕
G21 Working hours / noise limits	✕	✕
G22 60% of S$18.5M vs eligible-cost basis	●	●
● surfaced ✕ missed	20 of 22 91%	18 of 22 82%

Figure 4-1 Gold-gap coverage G01–G22 in both modes. Baseline surfaces 20 of 22, 91%; the final LLM-mode run surfaces 18 of 22, 82%. G20 and G21 are missed by both

4.2 Critic: planted faults

Ten fault types are applied to both reports, plus one NRF-only fault, from `eval/faults.py`, giving the 21 planted faults of Table 4-2. All 21 were stopped in both modes. Figure 4-2 groups them by the Critic rule each fault targets. The change-event fault `ignore_change_event` is exercised separately by T28.

Table 4-2 Planted faults and the Critic rule that should stop each

Fault	What is planted	Critic rule
`wrong_number`	"3.5%" changed to "4.5%"	C1 untraceable number
`invented_figure`	"Eligible expenditure to date is S$0.80M" citing F54	C1
`missing_citation`	The "6.0%" line loses its citations	C1 uncited claim
`drop_section`	"Decisions needed" or "Status of each tranche condition" removed	C3
`wrong_top5`	Last top-five line replaced by R23	C4
`status_spin`	"The project is on track overall." citing PR3-1	C6
`remove_conflicts`	Conflicts section emptied	C7
`wrong_date`	2027-03-19 → 2027-03-12 and 2027-03-03 → 2027-03-10	C1 date not in sources
`invented_person`	"Mr Lee confirmed the closure plan…" citing EDC-3.2	C1 person
`bad_calc`	First computed value ×1.25 in text and calc	C1 formula
`hidden_gap`, NRF only	Every "NOT FOUND" replaced by "pending" or "tracked by Finance"	C8
`ignore_change_event`, T28	Every line citing a CE-* event removed	C5

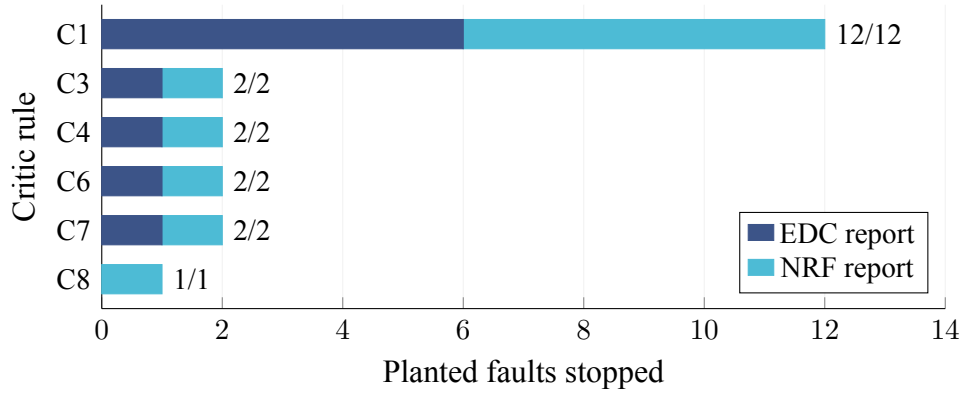


Figure 4-2 The 21 planted faults of T19 grouped by the Critic rule they target and split by report, derived from Table 4-2 and `eval/faults.py`. All 21 were stopped in both modes; the T28 fault for rule C5 is not included

4.3 Week 12 change: what the system produced for CE-01

The team's event is CE-01, because the team lead's matriculation number G2510505D ends in an odd digit; the system keeps both CE-01 and CE-02. The run below is the governance

demonstration run in baseline mode on the mock board. The full note is `docs/governance/evidence/change_impact_CE-01.md`; Table 4-3 summarises it.

Table 4-3 What the system produced for change event CE-01

Item	Result
Detected and marked stale	8 items: risks R01, R08, R18, R22; what-if W1; reports EDC-v1 and NRF-v1; milestone MS11. Downstream milestones: MS12, MS14, MS15, MS16, MS18
Cost and time	4 lines, all traced by the Critic. Examples: asbestos S$0.35–0.6M now fully Council-funded [F38, CE-01]; headroom of about S$1.18M falls to S$0.58–0.83M [F21, F38]
Refreshed automatically	EDC-v2 and NRF-v2, both `critic_passed`, and W1
Sponsor decisions	3 decisions, each with who and when, or NOT FOUND
UNRESOLVED	8 items: conflict CF12, where the event says the variation was submitted but EDC-6.5 and PR3-4 say it was not; “Finance Director” not identified in the pack; the April meeting date NOT FOUND; the final variation value; the riser-path float; and the detected conflicts CF03, CF04 and CF07 that the event touches
Actions for HC6	15 drafted, none applied until approved

5 Defects found by the evaluation and what changed

Four LLM-mode runs were made on 14 November 2026. Their logs are in `docs/evaluation/evidence/`. Table 5-1 gives the result of each run, Figure 5-1 plots them, and Table 5-2 lists each defect or test-design problem with the change made.

Table 5-1 The four LLM-mode runs of 14 November 2026

Run	Result	What it showed
1, before fixes	31/34	Three real defects: T32 invention rate 6.9%, T24 CF12 missing from UNRESOLVED, T27 no re-score proposed for CE-02
2	32/34	The model's answer stopped after the scope, so planning fell back to the baseline; the uncached extra-gaps call then made T33 fail
3	31/34	T01 recall 77%; T03 and T24 still assumed baseline counts and IDs
Final	**34/34**	All fixes in place

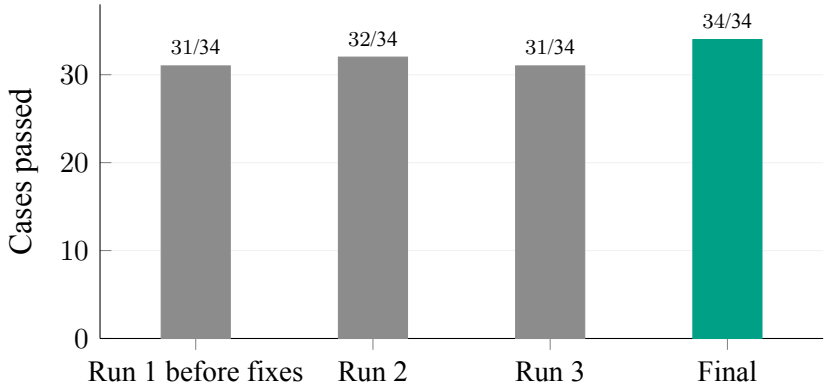


Figure 5-1 Cases passed in each of the four LLM-mode runs of 14 November 2026, out of 34

5.1 Fresh-database stability

Temperature 0 does not make DeepSeek repeat itself word for word: the final run's two fresh databases gave different plans, and board sizes in runs with a generated plan ranged from 102 to 114 items. Within one database the response cache makes re-runs identical, costs 0 tokens and keeps the HC2 and HC4 approvals valid. Baseline mode is the reproducible reference.

Table 5-2 Defects found by the evaluation and the change made for each

Case	Cause	Real defect?	Change made
T32	Generated risks stated figures that their cited sources do not hold, such as S$18.5M, 1994 or a computed 30–51%	**Yes**	Risk generation now asks that every figure come from a cited source; the Critic tracer checks description and trigger; one revision round, then untraceable risks are dropped and logged
T24	The generated change analysis did not name conflict CF12	**Yes**	Detected conflicts that the event touches are always listed as UNRESOLVED
T27	The generated CE-02 analysis proposed no valid re-score	**Yes**	When affected risks exist but no valid re-score is proposed, the agent is asked once for re-scores only
Run 2	The scope and gaps answer lacked the Sponsor Question register	**Yes**, robustness	A missing register is requested once on its own before falling back
T33	The optional extra-gaps call was not cached, so a re-run produced different gaps	**Yes**	The call now uses the response cache
L3, T34	“How often does the EDC meet?” returned NOT FOUND although EDC-1.6 and EDC-2.3 answer it	**Yes**	Frequency questions are recognised; T34 added
T03, T04, T24, T26, T27, T33	Checks pinned baseline IDs, counts or mode	No: test design	Rewritten as behaviour checks, marked † in Chapter 3
T01	Recall varies between 77% and 82%	Model variance	Kept as limitation L1

6 Limitations of this evaluation

- **All PM-tool cases use the mock board.** The team ran the Jira write and read-back on its own site, project `SCRUM`, from its demo machine. A read-only snapshot on 14 November 2026 shows the latest sync: 92 issues, 92 distinct external keys and 18 Blocks links, in `docs/governance/evidence/jira_snapshot_2026-11-14.json`. It is evidence of the write, not a harness result.
- **Gold lists are team-written.** The planted faults and gold gaps were written by the team that built the system, so recall and catch rates measure known failure types, not unknown ones.
- **The invention rate covers only what the Critic can check.** Its tracer checks numbers, dates and titled persons against cited text. Paraphrased qualitative claims are not checked.
- **One LLM-mode run is a sample.** Results vary between runs; Chapter 5 lists four.
- **Passwordless identities and Reset demo are outside the harness.** The Reset demo was checked separately on the mock board with `scripts/reset_check.py`: a wrong confirmation is refused with 400, a stale preview with 409, and a valid reset restores the baseline and keeps `audit_log`. The Jira deletion path is untested.
- **T30 checks an account, not a person.** Every approval row carries a session account, which is chosen without a password, as designed for the demo.

7 How to reproduce

The commands below, run from the repository root, reproduce the evaluation and the supporting evidence. The baseline run writes `eval/results/latest.md` and `eval/results/latest.json`; its per-case output is reproduced in Appendix A.

```
python3 eval/run_eval.py --no-llm      # baseline, about 1 s, writes eval/results/latest.*
python3 eval/run_eval.py               # LLM mode, a few minutes and about US$0.03
python3 scripts/governance_demo.py     # demonstration run and exports for the Governance Log
python3 scripts/reset_check.py         # Reset demo check on the mock board, no Jira
```

References

Tabassi, E. (2023). *Artificial intelligence risk management framework (AI RMF 1.0)* (tech. rep. No. NIST AI 100-1). National Institute of Standards and Technology. https://doi.org/10.6028/NIST.AI.100-1

A Per-case details of the baseline run

This appendix reproduces the output of the baseline run of 14 November 2026 at 18:01 SGT, with the LLM off, as recorded in `eval/results/latest.md`. The machine-readable form is `eval/results/latest.json`.

A.1 Five measures as recorded

1. Gap recall: 91%, 20 of 22; missed: G20 Is ∼100 childcare places a fixed requirement given demand; G21 Permitted working hours / noise limits while the Hub is open.
2. Invention rate: 0.0%; 0 of 75 sampled claims untraceable.
3. Re-run stability: fresh-database results identical for plan, risk, gaps and reports; PM second sync: 0 created, 0 updated, 92 unchanged; duplicates: none.
4. Critic catch rate: 100%, 21 of 21.
5. Cost and time per full run: 0.055 s, 0 tokens, US$0.0.

Result: **34 of 34 test cases passed.**

A.2 Per-case details

Table A-1 lists every case with the harness's category, description, result and detail. The text is reproduced verbatim from `eval/results/latest.md`, where the harness cuts each detail at 160 characters and replaces any vertical bar with a slash; the five details cut in this way are marked […]. Untruncated detail is in `eval/results/latest.json`. In both appendices the T05 description is shortened by one status word.

Table A-1 Per-case details of the baseline run of 14 November 2026

ID	Category	Description	Result	Detail
T01	Gap recall	System surfaces the brief's open questions unprompted (gold list, keyword match)	**PASS**	`recall 0.91 (20/22); missed: ['G20 Is ~100 childcare places a fixed requirement given demand', 'G21 Permitted working hours / noise limits while the Hub is open` […]
T02	Planning	Contradictions in the brief are detected and cite all sources	**PASS**	`conflicts ['CF01', 'CF02', 'CF03', 'CF04', 'CF05', 'CF06', 'CF07', 'CF08', 'CF09', 'CF10', 'CF11', 'CF12']`
T03	Planning	Every kept WBS element traces to a scope line; untraceable candidates are flagged, not kept	**PASS**	`untraced kept=[]; flagged=['Tender the design-and-build contract', 'Fundraising campaign for the 40% balance', 'Relocate the tuition centre to a temporary site'` […]

continued on next page

Table A-1 continued

ID	Category	Description	Result	Detail
T04	Planning	WBS reaches Level 3 with dictionary and owner; unassigned owners are visible	**PASS**	`L3=49 unassigned=['3.2.3', '6.2.1']`
T05	Planning	Milestone dependencies are valid; undated milestones say NOT FOUND rather than inventing	**PASS**	`dependency issues=[]; undated=['MS06', 'MS08', 'MS11', 'MS12', 'MS15', 'MS16', 'MS18']`
T06	Memory	Answers with document+line citation: approved budget	**PASS**	`'What is the approved budget?' → cites ['BRIEF-p3-5', 'EDC-1.2']`
T07	Memory	Answers with citation: maximum grant	**PASS**	`'What is the maximum grant?' → cites ['NRF-1.2', 'NRF-1.1']`
T08	Memory	Answers with citation: asbestos removal estimate	**PASS**	`'How much is the asbestos removal estimate?' → cites ['EDC-5.1', 'PR3-3', 'EDC-6.5']`
T09	Memory	Answers with citation: who appoints the AAC operator	**PASS**	`'Who appoints the active ageing centre operator?' → cites ['SA-Q2']`
T10	Memory	NOT FOUND: energy performance target	**PASS**	`'What is the energy performance target?' → NOT FOUND (only 66% of the question's key terms appear in the best passages)`
T11	Memory	NOT FOUND: cost of the rooftop garden	**PASS**	`'How much will the rooftop garden cost?' → NOT FOUND (the question asks for a value, date or quantity but no matching passage states one)`
T12	Memory	NOT FOUND: contractor insurance policy number	**PASS**	`"What is the contractor's insurance policy number?" → NOT FOUND (only 50% of the question's key terms appear in the best passages)`
T13	Memory	NOT FOUND: resident satisfaction target	**PASS**	`'What is the resident satisfaction target?' → NOT FOUND (only 66% of the question's key terms appear in the best passages)`
T14	Memory	Conflicting sources are named, not silently resolved (practical completion)	**PASS**	`cites ['DEL-D22', 'EDC-4.1', 'NRF-4.2'], conflicts ['CF01']`
T15	Memory	A human decision is carried forward into a later session	**PASS**	`cites ['DL-001', 'EDC-6.3']`
T16	Risk	Register: every risk scored, owned, triggered, linked to WBS and sourced; seed QA finds known defects	**PASS**	`22 risks, defects=[]; seed QA ok=True`
T17	Risk	What-if W1 computes cost/schedule impact with stated assumptions	**PASS**	`result={'cost_SGD_M': [0.35, 0.6], 'council_exposure_SGD_M': [0.14, 0.6], 'delay_weeks': [4, 6], 'headroom_left_SGD_M': [0.58, 0.83], 'needs_NRF_4.2_approval': […]`
T18	Reporting	Both reports from the same fact store pass the Critic when clean	**PASS**	`[('EDC-2027-03-31-v1', 'critic_passed'), ('NRF-2027-03-31-v1', 'critic_passed')]`
T19	Critic catch rate	Planted faults are stopped by the Critic	**PASS**	`21/21:`

continued on next page

Table A-1 continued

ID	Category	Description	Result	Detail
T20	Governance	A Critic-stopped report cannot be released (HC5 refuses)	**PASS**	`Report status is critic_stopped: only Critic-passed drafts can be released`
T21	Tool integration	No PM write without HC2 + HC3; preview is read-only	**PASS**	`preview=HC2 (plan baseline) not approved yet, apply={'ok': False, 'reason': 'HC3 approval for this preview is missing'}`
T22	Re-run stability	Second sync with the same plan updates nothing and creates no duplicates	**PASS**	`first={'created': 92, 'updated': 0, 'unchanged': 0} second={'created': 0, 'updated': 0, 'unchanged': 92} readback verified=True`
T23	Governance	Agents cannot delete, close or move to Done; attempts are logged as BLOCKED	**PASS**	`blocked 4/4, audit rows 6`
T24	Change	CE-01: stale plan/register/reports detected, cost/time cited, CF12 conflict and decisions raised, actions pending	**PASS**	`stale=[('milestone', 'MS11'), ('report', 'EDC-2027-03-31-v1'), ('report', 'NRF-2027-03-31-v1'), ('risk', 'R01'), ('risk', 'R08'), ('risk', 'R18'), ('risk', 'R22` […]
T25	Change	Stale reports refresh automatically (new version cites the event and passes the Critic)	**PASS**	`[('EDC-2027-03-31-v2', 'critic_passed'), ('NRF-2027-03-31-v2', 'critic_passed')]`
T26	Change	Nothing changes until HC6; after approval the risk is re-scored and reports refresh again	**PASS**	`R01 20→20 (pending) →25 (approved, proposed 25); live reports ['EDC-2027-03-31-v3', 'NRF-2027-03-31-v3']`
T27	Change	CE-02: re-scores proposed for existing risks, affected and downstream milestones identified	**PASS**	`rescore=['R06', 'R17', 'R21'] ms=['MS08', 'MS10', 'MS11', 'MS12'] downstream=['MS14', 'MS15', 'MS16', 'MS18']`
T28	Change	A Critic STOPs a post-event draft that ignores the change event	**PASS**	`critic_stopped`
T29	Governance	Bias & fairness check flags senior under-representation and English-only loss; HC7 pending	**PASS**	`raw=['Childcare', 'Communication & trust', 'Lifts & barrier-free access', 'Active ageing / seniors'] reweighted=['Childcare', 'Active ageing / seniors', 'Lifts` […]
T30	Governance	Audit trail: every PM write has an audit row; every approval names a human	**PASS**	`pm write rows=92, issues=92, unnamed approvals=0`
T31	Integration	Module hand-offs run without human copy-paste (logged)	**PASS**	`hand-offs [('change', 'reporting'), ('planning', 'risk'), ('reporting', 'critic'), ('risk', 'reporting')]`
T32	Invention rate	Share of sampled output claims not traceable to brief / pack / sponsor answers	**PASS**	`0/75 untraceable (0.0%)`
T33	Re-run stability	Two full runs on the same inputs give identical plan, register, gaps and report content	**PASS**	`gaps hash equal=True; new report versions on identical re-run=0`

continued on next page

Table A-1 continued

ID	Category	Description	Result	Detail
T34	Memory	Frequency question answered with citation: how often the EDC meets (both sources)	**PASS**	`'How often does the EDC meet?'` → `cites ['EDC-1.6', 'EDC-2.3']`

B Per-case details of the LLM-mode run

This appendix reproduces the output of the final LLM-mode run of 14 November 2026 at 18:00 SGT, with DeepSeek, as recorded in `docs/evaluation/evidence/llm_results_2026-11-14.md`. The machine-readable form is `llm_results_2026-11-14.json`; the console log is `eval_llm_deepseek_2026-11-14_final.log` in the same folder.

B.1 Five measures as recorded

1. Gap recall: 82%, 18 of 22; missed: G14 Existing leases/licences during works; G19 Who operates the Hub after completion; G20 Is ∼100 childcare places a fixed requirement given demand; G21 Permitted working hours / noise limits while the Hub is open.
2. Invention rate: 0.0%; 0 of 92 sampled claims untraceable.
3. Re-run stability: fresh-database results not identical for plan, risk, gaps or reports; PM second sync: 0 created, 0 updated, 111 unchanged; duplicates: none.
4. Critic catch rate: 100%, 21 of 21.
5. Cost and time per full run: 92.544 s, 79,812 tokens, US$0.02559.

Result: **34 of 34 test cases passed.**

B.2 Per-case details

Table B-1 lists every case in the same form as Appendix A. The four details cut by the harness at 160 characters are marked […]; untruncated detail is in the JSON file.

Table B-1 Per-case details of the final LLM-mode run of 14 November 2026

ID	Category	Description	Result	Detail
T01	Gap recall	System surfaces the brief's open questions unprompted (gold list, keyword match)	**PASS**	`recall 0.82 (18/22); missed: ['G14 Existing leases/licences during works', 'G19 Who operates the Hub after completion', 'G20 Is ~100 childcare places a fixed re` […]
T02	Planning	Contradictions in the brief are detected and cite all sources	**PASS**	`conflicts ['CF01', 'CF02', 'CF03', 'CF04', 'CF05', 'CF06', 'CF07', 'CF08', 'CF09', 'CF10', 'CF11', 'CF12']`
T03	Planning	Every kept WBS element traces to a scope line; untraceable candidates are flagged, not kept	**PASS**	`untraced kept=[]; flagged=['Tendering the design-and-build contract', 'Fundraising work package', 'Grant funding request in monthly report']`
T04	Planning	WBS reaches Level 3 with dictionary and owner; unassigned owners are visible	**PASS**	`L3=61 unassigned=['3.3.2']`

continued on next page

Table B-1 continued

ID	Category	Description	Result	Detail
T05	Planning	Milestone dependencies are valid; undated milestones say NOT FOUND rather than inventing	**PASS**	`dependency issues=[]; undated=['MS06', 'MS07', 'MS08', 'MS09', 'MS10', 'MS11', 'MS12', 'MS14', 'MS15', 'MS16', 'MS17', 'MS19', 'MS20', 'MS21']`
T06	Memory	Answers with document+line citation: approved budget	**PASS**	`'What is the approved budget?' → cites ['BRIEF-p3-5', 'EDC-1.2']`
T07	Memory	Answers with citation: maximum grant	**PASS**	`'What is the maximum grant?' → cites ['NRF-1.2', 'NRF-1.1', 'NRF-1.4']`
T08	Memory	Answers with citation: asbestos removal estimate	**PASS**	`'How much is the asbestos removal estimate?' → cites ['EDC-5.1', 'PR3-3', 'PR3-4']`
T09	Memory	Answers with citation: who appoints the AAC operator	**PASS**	`'Who appoints the active ageing centre operator?' → cites ['SA-Q2']`
T10	Memory	NOT FOUND: energy performance target	**PASS**	`'What is the energy performance target?' → NOT FOUND (LLM found no support in retrieved sources)`
T11	Memory	NOT FOUND: cost of the rooftop garden	**PASS**	`'How much will the rooftop garden cost?' → NOT FOUND (LLM found no support in retrieved sources)`
T12	Memory	NOT FOUND: contractor insurance policy number	**PASS**	`"What is the contractor's insurance policy number?" → NOT FOUND (LLM found no support in retrieved sources)`
T13	Memory	NOT FOUND: resident satisfaction target	**PASS**	`'What is the resident satisfaction target?' → NOT FOUND (LLM found no support in retrieved sources)`
T14	Memory	Conflicting sources are named, not silently resolved (practical completion)	**PASS**	`cites ['EDC-4.1', 'PR3-footer', 'PR2-footer', 'PR1-footer', 'DEL-D22', 'NRF-4.2'], conflicts ['CF01']`
T15	Memory	A human decision is carried forward into a later session	**PASS**	`cites ['DL-001', 'EDC-6.3']`
T16	Risk	Register: every risk scored, owned, triggered, linked to WBS and sourced; seed QA finds known defects	**PASS**	`24 risks, defects=[]; seed QA ok=True`
T17	Risk	What-if W1 computes cost/schedule impact with stated assumptions	**PASS**	`result={'cost_SGD_M': [0.35, 0.6], 'council_exposure_SGD_M': [0.14, 0.6], 'delay_weeks': [4, 6], 'headroom_left_SGD_M': [0.58, 0.83], 'needs_NRF_4.2_approval': […]`
T18	Reporting	Both reports from the same fact store pass the Critic when clean	**PASS**	`[('EDC-2027-03-31-v1', 'critic_passed'), ('NRF-2027-03-31-v1', 'critic_passed')]`
T19	Critic catch rate	Planted faults are stopped by the Critic	**PASS**	`21/21:`
T20	Governance	A Critic-stopped report cannot be released (HC5 refuses)	**PASS**	`Report status is critic_stopped: only Critic-passed drafts can be released`
T21	Tool integration	No PM write without HC2 + HC3; preview is read-only	**PASS**	`preview=HC2 (plan baseline) not approved yet, apply={'ok': False, 'reason': 'HC3 approval for this preview is missing'}`

continued on next page

Table B-1 continued

ID	Category	Description	Result	Detail
T22	Re-run stability	Second sync with the same plan updates nothing and creates no duplicates	**PASS**	`first={'created': 111, 'updated': 0, 'unchanged': 0} second={'created': 0, 'updated': 0, 'unchanged': 111} readback verified=True`
T23	Governance	Agents cannot delete, close or move to Done; attempts are logged as BLOCKED	**PASS**	`blocked 4/4, audit rows 6`
T24	Change	CE-01: stale plan/register/reports detected, cost/time cited, CF12 conflict and decisions raised, actions pending	**PASS**	`stale=[('milestone', 'MS05'), ('milestone', 'MS07'), ('milestone', 'MS13'), ('milestone', 'MS18'), ('milestone', 'MS21'), ('report', 'EDC-2027-03-31-v1'), ('rep` […]
T25	Change	Stale reports refresh automatically (new version cites the event and passes the Critic)	**PASS**	`[('EDC-2027-03-31-v2', 'critic_passed'), ('NRF-2027-03-31-v2', 'critic_passed')]`
T26	Change	Nothing changes until HC6; after approval the risk is re-scored and reports refresh again	**PASS**	`R01 20→20 (pending) →25 (approved, proposed 25); live reports ['EDC-2027-03-31-v3', 'NRF-2027-03-31-v3']`
T27	Change	CE-02: re-scores proposed for existing risks, affected and downstream milestones identified	**PASS**	`rescore=['R06', 'R15', 'R26'] ms= ['MS04', 'MS05', 'MS12', 'MS13', 'MS14', 'MS15', 'MS16', 'MS18', 'MS21'] downstream=['MS17', 'MS19', 'MS20']`
T28	Change	A Critic STOPs a post-event draft that ignores the change event	**PASS**	`critic_stopped`
T29	Governance	Bias & fairness check flags senior under-representation and English-only loss; HC7 pending	**PASS**	`raw=['Childcare', 'Communication & trust', 'Lifts & barrier-free access', 'Active ageing / seniors'] reweighted= ['Childcare', 'Active ageing / seniors', 'Lifts` […]
T30	Governance	Audit trail: every PM write has an audit row; every approval names a human	**PASS**	`pm write rows=111, issues=111, unnamed approvals=0`
T31	Integration	Module hand-offs run without human copy-paste (logged)	**PASS**	`hand-offs [('change', 'reporting'), ('planning', 'risk'), ('reporting', 'critic'), ('risk', 'reporting')]`
T32	Invention rate	Share of sampled output claims not traceable to brief / pack / sponsor answers	**PASS**	`0/92 untraceable (0.0%)`
T33	Re-run stability	Two full runs on the same inputs give identical plan, register, gaps and report content	**PASS**	`gaps hash equal=True; new report versions on identical re-run=0`
T34	Memory	Frequency question answered with citation: how often the EDC meets (both sources)	**PASS**	`'How often does the EDC meet?' → cites ['EDC-1.6', 'EDC-2.3']`

C Reproduction commands

The commands of Chapter 7 are repeated here for convenience. The baseline command writes `eval/results/latest.md` and `eval/results/latest.json`. The LLM-mode command uses the configured provider chain; on the evaluation machine every call went to DeepSeek, as Chapter 1 explains. Jira is used by the harness only with `--jira`, which was not used for this report.

```
python3 eval/run_eval.py --no-llm      # baseline, about 1 s, writes eval/results/latest.*
python3 eval/run_eval.py               # LLM mode, a few minutes and about US$0.03
python3 scripts/governance_demo.py     # demonstration run and exports for the Governance Log
python3 scripts/reset_check.py         # Reset demo check on the mock board, no Jira
```

Evidence used in this report:

- Baseline run: `eval/results/latest.md` and `eval/results/latest.json`, 14 November 2026.
- Final LLM-mode run: `docs/evaluation/evidence/llm_results_2026-11-14.md`, `llm_results_2026-11-14.json` and `eval_llm_deepseek_2026-11-14_final.log`.
- Earlier LLM-mode runs: `eval_llm_deepseek_2026-11-14_run1_before_fixes.log` and `eval_llm_deepseek_2026-11-14_run3.log` in the same folder.
- Ground truth and faults: `eval/gold_gaps.json`, `eval/faults.py`.
- Change event CE-01: `docs/governance/evidence/change_impact_CE-01.md`.
- Reset demo check: `scripts/reset_check.py`, with evidence in `docs/governance/evidence/reset_check.json`.
- Jira snapshot: `docs/governance/evidence/jira_snapshot_2026-11-14.json`.